

project1_Raya

June 16, 2021

```
[1]: import numpy as np
import pandas as pd
# for dimensionality reduction
from sklearn.decomposition import PCA
```

```
[2]: df_train = pd.read_csv('train.csv')
print('Size of training set: {} rows and {} columns'
      .format(*df_train.shape))
```

Size of training set: 4209 rows and 378 columns

```
[3]: df_train.head()
```

```
[3]:   ID      y  X0 X1  X2 X3 X4 X5 X6 X8 ... X375 X376 X377 X378 X379 \
0    0  130.81   k  v   at  a  d  u  j  o ...     0     0     1     0     0
1    6   88.53   k  t   av  e  d  y  l  o ...     1     0     0     0     0
2    7   76.26  az  w   n  c  d  x  j  x ...     0     0     0     0     0
3    9   80.62  az  t   n  f  d  x  l  e ...     0     0     0     0     0
4   13   78.02  az  v   n  f  d  h  d  n ...     0     0     0     0     0

      X380 X382 X383 X384 X385
0         0     0     0     0     0
1         0     0     0     0     0
2         0     1     0     0     0
3         0     0     0     0     0
4         0     0     0     0     0
```

[5 rows x 378 columns]

```
[4]: y_train = df_train['y'].values
```

```
[5]: cols = [c for c in df_train.columns if 'X' in c]
print('Number of features: {}'.format(len(cols)))

print('Feature types:')
df_train[cols].dtypes.value_counts()
```

Number of features: 376

Feature types:

```
[5]: int64      368
      object      8
      dtype: int64
```

```
[6]: counts = [[], [], []]
      for c in cols:
          typ = df_train[c].dtype
          uniq = len(np.unique(df_train[c]))
          if uniq == 1:
              counts[0].append(c)
          elif uniq == 2 and typ == np.int64:
              counts[1].append(c)
          else:
              counts[2].append(c)

      print('Constant features: {} Binary features: {} Categorical features: {}\n'
            .format(*[len(c) for c in counts]))
      print('Constant features:', counts[0])
      print('Categorical features:', counts[2])
```

Constant features: 12 Binary features: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289',
'X290', 'X293', 'X297', 'X330', 'X347']

Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']

```
[7]: df_test = pd.read_csv('test.csv')
```

```
[8]: # remove columns ID and Y from the data as they are not used for learning
      usable_columns = list(set(df_train.columns) - set(['ID', 'y']))
      y_train = df_train['y'].values
      id_test = df_test['ID'].values

      x_train = df_train[usable_columns]
      x_test = df_test[usable_columns]
```

```
[9]: #Check for null and unique values
      def check_missing_values(df):
          if df.isnull().any().any():
              print("There are missing values in the dataframe")
          else:
              print("There are no missing values in the dataframe")
      check_missing_values(x_train)
      check_missing_values(x_test)
```

There are no missing values in the dataframe
 There are no missing values in the dataframe

```
[10]: #If for any column(s), the variance is equal to zero, then you need to remove
      ↪ those variables
      #Apply label encoder
      for column in usable_columns:
          cardinality = len(np.unique(x_train[column]))
          if cardinality == 1:
              x_train.drop(column, axis=1) # Column with only one
              # value is useless so we drop it
              x_test.drop(column, axis=1)
          if cardinality > 2: # Column is categorical
              mapper = lambda x: sum([ord(digit) for digit in x])
              x_train[column] = x_train[column].apply(mapper)
              x_test[column] = x_test[column].apply(mapper)
      x_train.head()
```

/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:11:
 SettingWithCopyWarning:
 A value is trying to be set on a copy of a slice from a DataFrame.
 Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
# This is added back by InteractiveShellApp.init_path()
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:12:  

SettingWithCopyWarning:  

A value is trying to be set on a copy of a slice from a DataFrame.  

Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
if sys.path[0] == '':
```

```
[10]:
```

	X286	X350	X162	X190	X51	X277	X232	X312	X96	X257	...	X335	X288	\
0	0	0	0	0	0	0	0	0	0	0	...	0	0	
1	0	0	0	0	1	0	0	0	1	0	...	0	0	
2	1	1	1	0	1	0	1	0	1	0	...	0	0	
3	1	1	1	0	0	0	1	0	1	0	...	0	0	
4	1	1	1	0	1	0	1	0	1	0	...	0	0	

	X164	X252	X81	X291	X161	X293	X212	X150
0	0	0	0	0	0	0	0	1
1	0	0	0	0	0	0	0	1
2	0	1	0	0	0	0	0	1
3	0	1	0	1	0	0	0	1

4 0 1 0 0 0 0 0 1

[5 rows x 376 columns]

```
[11]: #Make sure the data is now changed into numericals
print('Feature types:')
x_train[cols].dtypes.value_counts()
```

Feature types:

```
[11]: int64      376
dtype: int64
```

```
[12]: #Perform dimensionality reduction
n_comp = 12
pca = PCA(n_components=n_comp, random_state=420)
pca2_results_train = pca.fit_transform(x_train)
pca2_results_test = pca.transform(x_test)
```

```
[13]: #Training using xgboost

import xgboost as xgb
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split

x_train, x_valid, y_train, y_valid = train_test_split(
    pca2_results_train,
    y_train, test_size=0.2,
    random_state=4242)

d_train = xgb.DMatrix(x_train, label=y_train)
d_valid = xgb.DMatrix(x_valid, label=y_valid)
d_test = xgb.DMatrix(x_test)
d_test = xgb.DMatrix(pca2_results_test)

params = {}
params['objective'] = 'reg:linear'
params['eta'] = 0.02
params['max_depth'] = 4

def xgb_r2_score(preds, dtrain):
    labels = dtrain.get_label()
    return 'r2', r2_score(labels, preds)

watchlist = [(d_train, 'train'), (d_valid, 'valid')]

clf = xgb.train(params, d_train,
```

```
1000, watchlist, early_stopping_rounds=50,  
feval=xgb_r2_score, maximize=True, verbose_eval=10)
```

[04:30:22] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linear is now deprecated in favor of reg:squarederror.

[0] train-rmse:99.14835 valid-rmse:98.26297 train-r2:-58.35295
valid-r2:-67.63754

Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 50 rounds.

[10] train-rmse:81.27653 valid-rmse:80.36433 train-r2:-38.88428
valid-r2:-44.91014

[20] train-rmse:66.71610 valid-rmse:65.77334 train-r2:-25.87403
valid-r2:-29.75260

[30] train-rmse:54.86957 valid-rmse:53.88974 train-r2:-17.17752
valid-r2:-19.64401

[40] train-rmse:45.24491 valid-rmse:44.21971 train-r2:-11.35979
valid-r2:-12.89996

[50] train-rmse:37.44729 valid-rmse:36.37237 train-r2:-7.46666
valid-r2:-8.40428

[60] train-rmse:31.14749 valid-rmse:30.01874 train-r2:-4.85757
valid-r2:-5.40571

[70] train-rmse:26.08662 valid-rmse:24.90890 train-r2:-3.10872
valid-r2:-3.41053

[80] train-rmse:22.04639 valid-rmse:20.83274 train-r2:-1.93458
valid-r2:-2.08514

[90] train-rmse:18.84403 valid-rmse:17.60281 train-r2:-1.14397
valid-r2:-1.20265

[100] train-rmse:16.33630 valid-rmse:15.08463 train-r2:-0.61131
valid-r2:-0.61753

[110] train-rmse:14.40370 valid-rmse:13.14924 train-r2:-0.25262
valid-r2:-0.22909

[120] train-rmse:12.92868 valid-rmse:11.69346 train-r2:-0.00921
valid-r2:0.02800

[130] train-rmse:11.80810 valid-rmse:10.62254 train-r2:0.15816
valid-r2:0.19788

[140] train-rmse:10.98601 valid-rmse:9.85892 train-r2:0.27130
valid-r2:0.30906

[150] train-rmse:10.37397 valid-rmse:9.33236 train-r2:0.35023
valid-r2:0.38089

[160] train-rmse:9.92029 valid-rmse:8.97081 train-r2:0.40582
valid-r2:0.42794

[170] train-rmse:9.59072 valid-rmse:8.72660 train-r2:0.44464
valid-r2:0.45866

[180] train-rmse:9.34406 valid-rmse:8.56794 train-r2:0.47284
valid-r2:0.47816

```

[190]   train-rmse:9.15732      valid-rmse:8.46753      train-r2:0.49370
valid-r2:0.49032
[200]   train-rmse:9.01338      valid-rmse:8.40309      train-r2:0.50949
valid-r2:0.49805
[210]   train-rmse:8.90767      valid-rmse:8.36135      train-r2:0.52093
valid-r2:0.50303
[220]   train-rmse:8.81830      valid-rmse:8.33586      train-r2:0.53049
valid-r2:0.50605
[230]   train-rmse:8.76528      valid-rmse:8.31968      train-r2:0.53612
valid-r2:0.50797
[240]   train-rmse:8.71715      valid-rmse:8.31027      train-r2:0.54120
valid-r2:0.50908
[250]   train-rmse:8.67918      valid-rmse:8.30535      train-r2:0.54519
valid-r2:0.50966
[260]   train-rmse:8.64532      valid-rmse:8.30037      train-r2:0.54873
valid-r2:0.51025
[270]   train-rmse:8.60983      valid-rmse:8.30106      train-r2:0.55243
valid-r2:0.51017
[280]   train-rmse:8.57386      valid-rmse:8.30142      train-r2:0.55616
valid-r2:0.51012
[290]   train-rmse:8.55184      valid-rmse:8.30257      train-r2:0.55844
valid-r2:0.50999
[300]   train-rmse:8.51677      valid-rmse:8.30632      train-r2:0.56205
valid-r2:0.50954
Stopping. Best iteration:
[259]   train-rmse:8.64905      valid-rmse:8.30028      train-r2:0.54834
valid-r2:0.51026

```

[14]: *#Predict your test_df values using xgboost*

```

p_test = clf.predict(d_test)

sub = pd.DataFrame()
sub['ID'] = id_test
sub['y'] = p_test
sub.to_csv('xgb.csv', index=False)

sub.head()

```

```

[14]:      ID      y
0     1  82.811882
1     2  97.109688
2     3  83.692238
3     4  77.155693
4     5 111.890274

```